

```

#include "opencv2/imgproc/imgproc.hpp"
#include "opencv2/highgui/highgui.hpp"
#include <cuda_runtime.h>
#include <stdio.h>
#include <stdlib.h>
#include <math.h>

// BUILD: nvcc sobelEdge.cu -o sobelEdge -lopencv_highgui -lopencv_core -lopencv_videoio

__global__ void rgb2gray(unsigned char *rgb, unsigned char *gray, int num) {
    // Get thread and block indexes
    int i = ((blockIdx.y * gridDim.x + blockIdx.x) * blockDim.y + threadIdx.y) * blockDim.x + threadIdx.x;
    int step = blockDim.x * blockDim.y * gridDim.x * gridDim.y;
    int j;
    int g;
    while(i < num) {
        j=3*i;
        g = 114*rgb[j] + 587*rgb[j+1] + 299*rgb[j+2];
        gray[i] = g/1000;
        i += step;
    }
}

#define Gray(u,v) gray[(u)*cols+(v)]

__global__ void sobelEdge(unsigned char *gray, unsigned char *edge, int rows, int cols) {
    int xpixel = threadIdx.x + blockDim.x * blockIdx.x;
    int ypixel = threadIdx.y + blockDim.y * blockIdx.y;
    int xstep = blockDim.x * gridDim.x;
    int ystep = blockDim.y * gridDim.y;
    float gx,gy;
    int nleftEdge,nrightEdge,ntopEdge,nbotEdge;
    unsigned char pixelVal;
    /* float s1[]={
        1, 0, -1,
        2, 0, -2,
        1, 0, -1
    };
    float s2[]={
        1, 2, 1,
        0, 0, 0,
        -1, -2, -1
    };*/
    while(xpixel<cols){
        while(ypixel<rows){
            int k,j;
            gx=gy=0.0f;
            //determind if the pixel is on any of the edges
            nleftEdge=(xpixel>0);
            nrightEdge=(xpixel<cols);
            ntopEdge=(ypixel>0);
            nbotEdge=(ypixel<rows);

            //calculate sobel derivative with edge cases
            if(nleftEdge) gx+=2*Gray(ypixel,xpixel-1);
            if(nleftEdge&&ntopEdge){
                pixelVal=Gray(ypixel-1,xpixel-1);
                gx+=pixelVal;
                gy+=pixelVal;
            }
            if(ntopEdge) gy+=2*Gray(ypixel-1,xpixel);
            if(nrightEdge&&ntopEdge){
                pixelVal=Gray(ypixel-1,xpixel+1);
                gx-=pixelVal;
                gy+=pixelVal;
            }
        }
    }
}

```

```

        if(nrightEdge) gx-=2*Gray(ypixel,xpixel+1);
        if(nrightEdge&&nbotEdge){
            pixelVal=Gray(ypixel+1,xpixel+1);
            gx-=pixelVal;
            gy-=pixelVal;
        }
        if(nbotEdge) gy-=2*Gray(ypixel+1,xpixel);
        if(nleftEdge&&nbotEdge){
            pixelVal=Gray(ypixel+1,xpixel-1);
            gx+=pixelVal;
            gy-=pixelVal;
        }
        //take the magnitude, sqrtf causing unknown issues.
        edge[ypixel*cols+xpixel]=(unsigned char) abs(gx)+abs(gy);
        ypixel+=ystep;
    }
    ypixel=threadIdx.y+blockDim.y*blockIdx.y;
    xpixel+=xstep;
}

}

int main(int argc, char *argv[]) {

    // OpenCV stuff: camera, window, frame, etc.

    cv::VideoCapture cam("nvcamerasrc ! video/x-raw(memory:NVMM), width=(int)640, height=(int)480,format=(string)I420, framerate=(fraction)30/1 ! nvvidconv ! video/x-raw, format=(string)BGRx ! videoconvert ! video/x-raw, format=(string)BGR ! appsink");
    if(cam.isOpened()) { printf("camera opened\n"); }
    else { printf("camera not opened\n"); return 1; }
    int width = cam.get(CV_CAP_PROP_FRAME_WIDTH );
    int height = cam.get(CV_CAP_PROP_FRAME_HEIGHT);
    printf("width = %d, height = %d\n",width,height);
    int frame_size = width*height*3;
    int gray_size = width*height;
    unsigned char *h_frame_buff = (unsigned char*)calloc(frame_size,sizeof(unsigned char));
    unsigned char *h_gray_buff = (unsigned char*)calloc( gray_size,sizeof(unsigned char));
    unsigned char *h_sobel_buff = (unsigned char*)calloc( gray_size,sizeof(unsigned char));
    cv::Mat frame(height,width,CV_8UC3,h_frame_buff);
    cv::Mat gray(height,width,CV_8UC1, h_gray_buff);
    cv::Mat sobel(height,width,CV_8UC1, h_sobel_buff);
    cv::namedWindow( "Frame", 0 );
    cv::namedWindow( "Gray" , 1 );
    cv::namedWindow( "Sobel", 1 );

    // Allocate global memory on device
    unsigned char *d_frame_buff;
    unsigned char *d_gray_buff;
    unsigned char *d_sobel_buff;
    cudaMalloc((void**)&d_frame_buff,frame_size*sizeof(unsigned char));
    cudaMalloc((void**)&d_gray_buff , gray_size*sizeof(unsigned char));
    cudaMalloc((void**)&d_sobel_buff, gray_size*sizeof(unsigned char));

    // Define block and thread dimensions
    dim3 grid_size_rgb(16,16);
    dim3 block_size_rgb(16,16);

    cudaEvent_t start,stop;
    cudaEventCreate(&start);
    cudaEventCreate(&stop);
    float averageTime=0;
    // Processing loop
    int i;
    for(i=0;i<100;i++) { // For profiling

```

```

//for(;;) { // Run in an infinite loop
// Read in a camera frame
cam >> frame;
if( frame.empty() ) { printf("frame empty\n"); break; }

// Malloc frame to device
cudaMemcpy(d_frame_buff,h_frame_buff,frame_size*sizeof(unsigned char),cudaMemcpyHostToDevice);

cudaEventRecord(start);
// Launch kernels
rgb2gray<<<grid_size_rgb,block_size_rgb>>>(d_frame_buff,d_gray_buff,gray_size);
sobelEdge<<<grid_size_rgb,block_size_rgb>>>(d_gray_buff,d_sobel_buff,height,width);
cudaEventRecord(stop);

// Malloc result to host
cudaMemcpy(h_gray_buff,d_gray_buff,gray_size*sizeof(unsigned char),cudaMemcpyDeviceToHost);
cudaMemcpy(h_sobel_buff,d_sobel_buff,gray_size*sizeof(unsigned char),cudaMemcpyDeviceToHost);

cudaEventSynchronize(stop);
float milliseconds=0;
cudaEventElapsedTime(&milliseconds, start, stop);
averageTime+=milliseconds;
//printf("Elapsed time %f ms\n", milliseconds);

// Show the results
cv::imshow( "Frame", frame );
cv::imshow( "Gray", gray );
cv::imshow( "Sobel", sobel );
cv::waitKey(1);
}
averageTime/=100;
printf("Average Elapsed time %f ms\n", averageTime);
printf("Required time for 30 fps: %f ms\n",100.0f/30.0f);

cudaFree(d_frame_buff);
cudaFree(d_gray_buff);
cudaFree(d_sobel_buff);

// For profiling
cudaDeviceReset();

return 0;
}

```